

How to Train Your Intercept

A Comparison of Intercept Update Strategies in Coordinate Descent

Johan Larsson ¹ Department of Mathematical Sciences, University of Copenhagen

Date published: 2025-07-14 Last modified: 2025-07-14

Abstract

Training intercepts in regularized regression is a common problem, yet it is often overlooked in the literature. This document explores different strategies for updating intercepts in coordinate descent algorithms, comparing gradient, Newton, and exact strategies.

Keywords: intercepts, coordinate descent, optimization, regularization

Contents

1	1	Introduction	1
3		References	1

4 TaskLocalRNG()

1 Introduction

Intercept updating is a overlooked issue in the field of statistical learning. For regularized regression, coordinate descent is the dominating algorithm, yet it is not clear how to best update the intercept term. Implementations differ in how they handle this, for instance `skglm` uses a gradient step, `LIBLINEAR` uses a Newton step, and `blitz` uses an exact update (based on Newton steps).

References

¹Corresponding author: jola@math.ku.dk




Figure 1: Parametric Plots